

Stats 326: Assignment 3

Hasnain Cheena

08/05/2020

Question 1

```
#seasonal factor model
sf.fit = lm(CO2.fit.ts[-1] ~ reduced.Time[-1] + reduced.Time.break[-1] + reduced.Quarter[-1] + CO2.fit.ts[-75])

#model summary
summary(sf.fit)

##
## Call:
## lm(formula = CO2.fit.ts[-1] ~ reduced.Time[-1] + reduced.Time.break[-1] +
##     reduced.Quarter[-1] + CO2.fit.ts[-75])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.48990 -0.12209 -0.00581  0.11306  0.53995
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    108.82449    29.39775   3.702 0.000435 ***
## reduced.Time[-1]    0.14548     0.03860   3.769 0.000349 ***
## reduced.Time.break[-1] 0.04182     0.01196   3.496 0.000843 ***
## reduced.Quarter[-1]2  0.43876     0.07593   5.778 2.14e-07 ***
## reduced.Quarter[-1]3  1.14763     0.07477  15.348 < 2e-16 ***
## reduced.Quarter[-1]4  0.41510     0.06543   6.344 2.22e-08 ***
## CO2.fit.ts[-75]      0.70187     0.08043   8.727 1.18e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1889 on 67 degrees of freedom
## Multiple R-squared:  0.9997, Adjusted R-squared:  0.9997
## F-statistic: 4.282e+04 on 6 and 67 DF, p-value: < 2.2e-16

#forecast 2018 Q4
t76.sf.pred = sf.fit$coefficients[1] + (sf.fit$coefficients[2] * 76) + (sf.fit$coefficients[3] * 26) +
  (sf.fit$coefficients[6]*1) + (sf.fit$coefficients[7]*CO2.fit.ts[75])

#forecast 2019 Q1
t77.sf.pred = sf.fit$coefficients[1] + (sf.fit$coefficients[2] * 77) + (sf.fit$coefficients[3] * 27) +
  (sf.fit$coefficients[7]*t76.sf.pred)

#forecast 2019 Q2
t78.sf.pred = sf.fit$coefficients[1] + (sf.fit$coefficients[2] * 78) + (sf.fit$coefficients[3] * 28) +
  (sf.fit$coefficients[7]*t77.sf.pred)
```

```

icients[3] * 28) +
  (sf.fit$coefficients[4]*1) + (sf.fit$coefficients[7]*t77.sf.pred)

#forecast 2019 Q3
t79.sf.pred = sf.fit$coefficients[1] + (sf.fit$coefficients[2] * 79) + (sf.fit$coeffi
icients[3] * 29) +
  (sf.fit$coefficients[5]*1) + (sf.fit$coefficients[7]*t78.sf.pred)

#results
results.sf.df = data.frame(Time=c("2018.4", "2019.1", "2019.2", "2019.3"),
                           Predictions=c(t76.sf.pred,t77.sf.pred,t78.sf.pred,
t79.sf.pred))
results.sf.df

##      Time Predictions
## 1 2018.4      406.0347
## 2 2019.1      406.1401
## 3 2019.2      406.8401
## 4 2019.3      408.2276

#RMSEP
sf.RMSEP = sqrt(1/4*sum((results.sf.df$Predictions-CO2.pred.ts)^2))
sf.RMSEP

## [1] 0.2384888

```

The Seasonal Factor model included a Time variable, break-in trend variable, seasonal factor and lagged response variable.

The Residual Series showed reasonably constant scatter about 0. There is a slight upward trend in the early part of the Residual Series. Furthermore, there is a large negative residual at time period 38 (2009.2) and a large positive residual at time period 66 (2016.2).

The plot of the autocorrelation function shows weakly significant lags at 1, 11 and 16.

The residuals appear to follow a normal distribution (Shapiro-Wilk p-value = 0.85, min = -0.49 and max=0.54).

The concentrations of Quarters 2 to 4 were all higher than the omitted baseline level (Quarter 1) with Quarter 3 being the highest (1.15ppm above Quarter 1).

The RMSEP indicates that, on average, the prediction error is 0.24 ppm.

Question 2

```
#model fit
FH.fit = lm(CO2.fit.ts[-1] ~ reduced.Time[-1] + reduced.Time.break[-1] + c1[-1] + s1[-1] + c2[-1] + CO2.fit.ts[-75])
#model summary
summary(FH.fit)

##
## Call:
## lm(formula = CO2.fit.ts[-1] ~ reduced.Time[-1] + reduced.Time.break[-1] +
##      c1[-1] + s1[-1] + c2[-1] + CO2.fit.ts[-75])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.48990 -0.12209 -0.00581  0.11306  0.53995
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    109.32486    29.38088   3.721 0.000408 ***
## reduced.Time[-1]    0.14548     0.03860   3.769 0.000349 ***
## reduced.Time.break[-1]  0.04182     0.01196   3.496 0.000843 ***
## c1[-1]          -0.01183     0.04370  -0.271 0.787404
## s1[-1]          -0.57381     0.03739 -15.348 < 2e-16 ***
## c2[-1]          -0.07344     0.02232  -3.290 0.001597 **
## CO2.fit.ts[-75]     0.70187     0.08043   8.727 1.18e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1889 on 67 degrees of freedom
## Multiple R-squared:  0.9997, Adjusted R-squared:  0.9997
## F-statistic: 4.282e+04 on 6 and 67 DF, p-value: < 2.2e-16

#forecast 2018 Q4
t76.FH.pred = FH.fit$coefficients[1] + (FH.fit$coefficients[2] * 76) + (FH.fit$coefficients[3] * 26) +
(FH.fit$coefficients[4] * cos(2*pi*76*(1/4))) +
(FH.fit$coefficients[5] * sin(2*pi*76*(1/4))) +
(FH.fit$coefficients[6] * cos(2*pi*76*(2/4))) +
(FH.fit$coefficients[7]*CO2.fit.ts[75])

#forecast 2019 Q1
t77.FH.pred = FH.fit$coefficients[1] + (FH.fit$coefficients[2] * 77) + (FH.fit$coefficients[3] * 27) +
(FH.fit$coefficients[4] * cos(2*pi*77*(1/4))) +
(FH.fit$coefficients[5] * sin(2*pi*77*(1/4))) +
(FH.fit$coefficients[6] * cos(2*pi*77*(2/4))) +
(FH.fit$coefficients[7]*t76.FH.pred)
```

```

#forecast 2019 Q2
t78.FH.pred = FH.fit$coefficients[1] + (FH.fit$coefficients[2] * 78) + (FH.fit$coefficients[3] * 28) +
(FH.fit$coefficients[4] *cos(2*pi*78*(1/4))) +
(FH.fit$coefficients[5] *sin(2*pi*78*(1/4))) +
(FH.fit$coefficients[6] *cos(2*pi*78*(2/4))) +
(FH.fit$coefficients[7]*t77.FH.pred)

#forecast 2019 Q3
t79.FH.pred = FH.fit$coefficients[1] + (FH.fit$coefficients[2] * 79) + (FH.fit$coefficients[3] * 29) +
(FH.fit$coefficients[4] *cos(2*pi*79*(1/4))) +
(FH.fit$coefficients[5] *sin(2*pi*79*(1/4))) +
(FH.fit$coefficients[6] *cos(2*pi*79*(2/4))) +
(FH.fit$coefficients[7]*t78.FH.pred)

#results
results.FH.df = data.frame(Time=c("2018.4", "2019.1", "2019.2", "2019.3"),
                             Predictions=c(t76.FH.pred,t77.FH.pred,t78.FH.pred,
t79.FH.pred))
results.FH.df

##      Time Predictions
## 1 2018.4      406.0347
## 2 2019.1      406.1401
## 3 2019.2      406.8401
## 4 2019.3      408.2276

#RMSEP
FH.RMSEP = sqrt(1/4*sum((results.FH.df$Predictions-CO2.pred.ts)^2))
FH.RMSEP

## [1] 0.2384888

```

The Full Harmonic model produced the same results as the Seasonal Factor, which is expected. The Full Harmonic model had the smallest RMSEP of all the Harmonic models (0.238 ppm).

The Full Harmonic model included a Time variable, break-in trend variable, 3 harmonics (c1 with a P-value = 0.79 being non-significant) and a lagged response variable. The Residual Series showed reasonably constant scatter about 0. There is a slight upward trend in the early part of the Residual Series. Furthermore, there is a large negative residual at time period 38 (2009.2) and a large positive residual at time period 66 (2016.2). The plot of the autocorrelation function shows weakly significant lags at 1, 11 and 16. The residuals appear to follow a normal distribution (Shapiro-Wilk p-value = 0.85, min=-0.49 and max=0.54).

The reduced harmonic model created by removing non-significant pairs of harmonics produced the same model as Full Harmonic model (because there were no non-significant pairs). Another reduced harmonic model created by removing non-significant harmonics resulted in removing c1 and produced an RMSEP of 0.243 ppm. Therefore, it was rejected.

Furthermore, a cosine model was not appropriate as the observations did not follow a smooth harmonic curve for each year.

Question 3

The Seasonal Factor model included a Time variable, break-in trend variable, seasonal factor and lagged response variable to account for the autocorrelation.

The Residual Series showed reasonably constant scatter about 0. There is a slight upward trend in the beginning of the Residual Series. However, as this is in the early part of the series it is not of concern. There is a large negative residual for time period 38 (2009.2) and a large positive residual for time period 66 (2016.2). The plot of the autocorrelation function shows weakly significant autocorrelation at lag 1, lag 11 and lag 16. As these lags are weakly significant, they are not a worry. The residuals appeared to follow a normal distribution (Shapiro-Wilk P-value = 0.852, min=-0.49 and max=0.54). Therefore, model assumptions are satisfied.

We have strong evidence against the hypothesis that the coefficient associated time variable is 0 (p-value = 0.0003). Further, we have strong evidence against the hypothesis that the coefficient associated with the break-in trend time variable is 0 (p-value = 0.0008). Additionally, we have very strong evidence against the hypothesis of no autocorrelation (p-value = 0).

The F-statistic provides extremely strong evidence against the hypothesis that none of the variables are related to the CO2 concentration (p-value 0). The Multiple R^2 is almost 1 (0.9997) indicating that nearly all the variation in the CO2 concentration is explained by the model.

The Residual Standard Error is 0.19 ppm so the prediction intervals should be reasonably narrow. The model predictions can be relied upon as the assumptions are satisfied. The RMSEP for the predictions from 2018 Q4 to 2019 Q3 was 0.238 ppm which was smaller than the Reduced Harmonic model (0.244 ppm) and the same as the Full Harmonic model.

The predictions for 2018 Q4 to 2019 Q3 are:

2018 Q4: 406.03 ppm
2019 Q1: 406.14 ppm
2019 Q2: 406.84 ppm
2019 Q3: 408.23 ppm

Question 4

```
#model fit
sf.full.fit = lm(CO2.full.ts[-1] ~ Time[-1] + Time.break[-1]+ Quarter[-1] + CO2.full.
ts[-79])
summary(sf.full.fit)

##
## Call:
## lm(formula = CO2.full.ts[-1] ~ Time[-1] + Time.break[-1] + Quarter[-1] +
##     CO2.full.ts[-79])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.50285 -0.12867 -0.00066  0.12102  0.53787
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    108.36491    28.78665   3.764 0.000341 ***
## Time[-1]         0.14500     0.03785   3.831 0.000273 ***
## Time.break[-1]   0.04082     0.01114   3.665 0.000474 ***
## Quarter[-1]2     0.46150     0.07438   6.205 3.25e-08 ***
## Quarter[-1]3     1.16834     0.07242  16.132 < 2e-16 ***
## Quarter[-1]4     0.41786     0.06380   6.549 7.79e-09 ***
## CO2.full.ts[-79]  0.70309     0.07876   8.927 3.20e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1888 on 71 degrees of freedom
## Multiple R-squared:  0.9998, Adjusted R-squared:  0.9997
## F-statistic: 5.127e+04 on 6 and 71 DF,  p-value: < 2.2e-16

#forecast 2019 Q4
t80.sf.pred = sf.full.fit$coefficients[1] + (sf.full.fit$coefficients[2] * 80) + (sf.
full.fit$coefficients[3] * 30) +
  (sf.full.fit$coefficients[6]*1) + (sf.full.fit$coefficients[7]*CO2.full.ts[79])

#forecast 2020 Q1
t81.sf.pred = sf.full.fit$coefficients[1] + (sf.full.fit$coefficients[2] * 81) + (sf.
full.fit$coefficients[3] * 31) +
  (sf.full.fit$coefficients[7]*t80.sf.pred)

#forecast 2020 Q2
t82.sf.pred = sf.full.fit$coefficients[1] + (sf.full.fit$coefficients[2] * 82) + (sf.
full.fit$coefficients[3] * 32) +
  (sf.full.fit$coefficients[4]*1) + (sf.full.fit$coefficients[7]*t81.sf.pred)

#forecast 2020 Q3
t83.sf.pred = sf.full.fit$coefficients[1] + (sf.full.fit$coefficients[2] * 83) + (sf.
full.fit$coefficients[3] * 33) +
  (sf.full.fit$coefficients[5]*1) + (sf.full.fit$coefficients[7]*t82.sf.pred)
```

```
#results
results.sf.df = data.frame(Time=c("2019.4", "2020.1", "2020.2", "2020.3"),
                             Predictions=c(t80.sf.pred,t81.sf.pred,t82.sf.pred,
t83.sf.pred))
results.sf.df

##      Time Predictions
## 1 2019.4      408.6447
## 2 2020.1      408.6901
## 3 2020.2      409.3694
## 4 2020.3      410.7396
```

The summary shows the seasonal factor model parameter estimates for the full timeframe model (*sf.full.fit*) are similar to those of the reduced timeframe model (*sf.fit*). As the model assumptions were satisfied our predictions should be reasonably reliable.

The prediction intervals are between 0.74 and 0.98 ppm.

Question 5

The best predicting model is the seasonal-trend-lowess seasonally adjusted model as it had the lowest RMSEP (0.2 ppm). As the model assumptions for the seasonally adjusted model were satisfied, we should be able to rely on any predictions.

The prediction intervals for the Seasonally Adjusted and Seasonal Factor models are the same (both are 0.74 to 0.98 ppm).